

为读取共享数据的客户端提供更多类型的数据一致性保证措施是可行的，可能也是必须的。

作者：DOUG TERRY

以棒球赛为例 解释复制数据 的一致性问題

云中的复制存储系统为读取数据的应用程序提供了各种不同的一致性保证措施。云存储供应商总是以冗余方式将数据存储在台机器上，从而确保在发生不可避免的故障时，数据仍然可用。跨数据中心进行数据复制的情况并不少见，使得数据在某个数据中心完全瘫痪的情况下有可能幸免于难。不过，这些数据副本并不是总能保持完全同步。因此，客户端从不同的服务器读取同一数据对象时，有可能接收到不同的数据版本。

某些系统，如微软的 Windows Azure，对于应用程序只提供强一致性存储服务。⁵ 这些服务确保 Windows Azure 存储的客户端看到的总是数据对象的最新值。尽管在数据中心内部提供强一致性存储服务是一种符合用户期待的合理选择，但当系统开始为位于多个大洲的多个数据中心提供跨越地理范围的数据复制服务时，却带来了问题。

许多云存储系统，如亚马逊的“简单存储服务”（S3），在设计时采用了弱一致性。其设计者认为强一致性对于大型系统而言代价过大。他们选择松弛一致性，以获得更好的性能和可用性。客户端在此类系统中执行读操作时，可能会得到过期的数据。读操作所返回的数据，是过去某个时间点上该对象的值，但不一定是最新的值。例如，当读操作指向了一个副本，而该副本尚未像别的副本一样收到所有写操作时，就会发生这种情况。此类系统被称作支持最终一致性。¹²

近来设计的系统都认识到支持不同类型应用程序的必要性，提供了访问云存储的多种操作方式。例如，亚马逊的 DynamoDB 同时提供了最终一致性读取方式和强一致性读取方式，后者需要较长的读取等待时间，并使读取吞吐量下降一半。¹ 亚马逊 SimpleDB 系统为读取数据的客户端提供了同样的选择。与此类似，Google App Engine Datastore（谷歌应用引擎数据存储）系统新增了最终一致性读取方式，对其默认的强一致性读取方式进行补充。⁸ 雅虎公司许多 Web 服务底层所使用的 PNUTS，提供了三种类型的读操作：任意读（read-any）、临界读（read-critical）

» 重要见解

- 虽然复制云服务通常提供强一致性或最终一致性，但位于两者之间的一致性保证可能会更好地满足应用程序的需求。
- 一致性保证可以使用一种与实现方案无关的方式来定义，而且可以在每次进行读操作时加以选择。
- 对于松弛一致性的处理，不一定会给应用程序开发人员带来额外的负担。



和最新读 (*read-latest*)。 ⁷ 基于配额的新型存储系统允许客户端通过选择不同的读写配额来选择强一致性或最终一致性。 ⁴

研究界在过去 30 年中为分布式复制系统提出了若干数据一致性模型。 ¹⁰ 这些模型提供介于强一致性和最终一致性两者之间的一致性保证。 例如， 某个系统可能会

保证客户端看到的数据不会是 5 分钟之前的过时数据， 或保证客户端始终看到自己写入的结果。 实际上， 一些一致性模型甚至比最终一致性模型还要弱。 但鉴于这些模型用处不大， 本文略去对它们的讨论。

之所以要探索不同的一致性模型， 是因为一致性、 性能和

可用性 (简称 CAP) 不能同时达到最优。 ^{9,10,12,13} 提供更强的一致性， 通常会造成较低的读写性能和较差的数据可用性。 CAP 定理已经证明， 对于必须容忍网络不连通的系统， 设计人员必须在一致性和可用性之间进行抉择。 ⁵ 在实践中， 访问延迟也是同样重要的一个考虑因素。 ¹ 在各种复杂的利弊

表 1. 六种一致性保证措施。

强一致性	看到之前的所有写入。
最终一致性	看到之前写入的子集。
一致性前缀	看到初始写入序列。
限定过期法	看到所有“旧”的写入。
单调读	看到写入的增量子集。
读自身写	看到读取者过去的所有写操作。

权衡下，人们提出的每一种一致性模型都有某种意义。

不同的一致性有实际用处吗？应用程序开发人员能应对最终一致性吗？现有的云存储服务已经提供了一致性读方式和最终一致性读方式，系统是否应提供更多的一致性选项呢？

通过研究棒球赛这一应用示例（显然是虚构示例），本文试图回答这些问题，至少是部分回答这些问题。需要了解棒球赛得分的人包括记分员、裁判员、电台记者、体育新闻记者和统计员，下面我将详细探讨不同人员的各种需求。假设比赛得分存储在一个基于云的复制存储服务中，我要说明最终一致性对于大多数参与者来讲是不够的，而强一致性也是不必要的。大多数参与者会受益于一些中间的一致性。

下一节将定义针对读操作的六种可能的一致性保证。然后，我将给出一种模拟棒球赛的算法，其中表明写入和读取数据是在何处发生的。当采用不同的一致性保证措施读取比分时，我会将所有可能返回的结果——列举出来。我还将研究想知道棒球比分数据的这些人所担当的不同角色，以及每种角色所期望的读一致性，并从这个简单的例子中得出我的结论。

读一致性保证

尽管在过去 30 年里，复制系统提供了多种数据一致性措施，计算机科学研究界也已探索出各种各样的一致性模型，但是这些模型

很多都限定在特定的实现方案中。要理解系统在什么情况下能够提供什么样的一致性，经常需要了解系统是如何运作的。不幸的是，这把负担加到那些为此类存储系统开发应用程序的人身上。

笔者此处提倡的六种一致性保证可以用一种简单的、与实现无关的方式进行描述。这不仅能让应用程序开发人员获益，还能给底层存储系统的设计、运行和演化带来灵活性。

这些一致性保证是基于客户端要对数据存储区进行读写操作这样一个简单模型。多个客户端可能同时访问共享的信息，如社会关系图、新闻推送、照片、购物车或财务记录。数据在一组服务器间复制，而复制协议的细节对客户端是不可见的。写操作是指对一个或多个数据对象进行更新的任何操作。所有服务器最终都会收到这些写操作，并以相同的顺序加以执行。该顺序与客户端提交写操作的顺序是一致的。实际上，可以通过在主服务器上执行所有这些写操作来保证这个顺序，或者通过在每个服务器上运行一种一致性协议来就全局顺序达成一致。读操作返回先前写入一个或多个数据对象的值，但不一定是最新的值。每个读操作可以请求某种一致性保证，所选择的一致性决定了允许返回值集合。每种一致性保证由哪些先前写入的结果对读操作可见定义。表 1 总结了六种一致性保证

强一致性特别容易理解。这种一致性确保读操作返回的值就是

最后一次写入给定对象的值。如果写操作可以修改或扩充部分数据对象，例如将数据追加到日志的末尾，那么，读操作会返回将所有写操作应用到该对象后的结果。换句话说，读操作看到的是所有先前写操作的结果。

最终一致性是最弱的一种一致性保证。这意味着它的可能返回值集合最为庞大。对于整个对象的写入，最终一致性读操作可能返回曾经写入该数据对象的任何值。广而言之，此读操作可能会从一个副本返回结果，该副本可能收到了对被读取数据对象写操作的任意子集。术语“最终”一致性源于这样一个事实：每个副本最终会收到每一个写操作。如果客户端停止了写操作，那么读操作最终将返回对象的最新值。

通过请求一致性前缀，可以保证数据读取者能够看到从第一次写入数据存储区开始的一个有序的写入结果序列。例如，读操作可能从一个副本得到应答，而该副本从主副本那里依次接收写操作，但尚未收到最近的一些写操作。换句话说，数据读取者看到的是数据存储区的某个版本，而该版本是在过去某个时间存在于主副本上的。这类似于许多数据库管理系统所提供的“快照隔离”一致性。对于读取单个数据对象的操作，如果系统中写操作完全覆盖该对象以前的值，那么即便是最终一致性的读操作也能看到一致性前缀。当读取多个数据对象，或写操作增量更新一个对象时，就能体现出请求一致性前缀的好处。

限定过期法确保读取的结果不会太陈旧。典型情况下，用时间段 T 定义过期，比如五分钟。存储系统确保读操作返回不超过 T 分钟以前写入的值或是更晚写入的值。还有一些系统采用另外的方式来定义数据过期：缺失的写入次数，甚至是数据值的不准确

程度。我觉得时间限定过期法对应用程序开发人员来说是最自然的概念。

单调读 (monotonic read) 是一种属性，适用于由某一存储系统客户端执行的读操作序列。因此称之为“会话期保证”。¹¹ 对于单调读，客户端可以读取任意陈旧数据，这一点与最终一致性相同，但可以保证所观察到的数据存储区随着时间的推移会越来越新。具体而言，如果客户端发出一个读操作，稍后再发出针对同一对象的另一个读操作，那么，第二次读操作返回的将是相同的值或更晚写入的值。

读自身写 (read my writes) 是一种属性，也适用于由单个客户端执行的一系列操作。它保证由客户端完成的所有写操作的结果，对该客户端的后续读操作而言都是可见的。如果客户端向一个数据对象写入一个新值，然后再读取这个对象，那么，读操作将会返回该客户端最后写入的值（或者是由另一个不同的客户端在此之后写入的其他值）。对于没有发出过写操作的客户端来说，这种保证与最终一致性保证相同。（注：以前的文章中称之为“读取你的写入”，¹¹ 不过为了更准确地描述这种一致性保证，我对其从客户端的角度进行了重新命名。）

最后这四种读保证都是最终一致性的一种形式，但要比云存储系统通常提供的最终一致性模型更强。每种一致性保证的“强度”，不依赖于数据何时以及如何和服务端间传播，而是取决于一个读操作所允许的结果集大小。较小的读结果集表明较强的一致性。当要求强一致性时，只能返回一个结果，即已被写入的最新值。对于已经更新过多次的对象，最终一致性读操作可能会返回众多适合值中的一个。所有四个中间级别的一致性保证，彼此没有强弱之

分。也就是说，对于某个读操作，每种保证可能会有一组不同的响应。在某些情况下，应用程序可能需要请求多种保证，如稍后所示。例如，客户端可以同时请求单调读和读自身写两种保证，以便它看到的数据存储区与其自身进行的操作保持一致。¹¹

本文用于记录棒球得分的数据存储区是一个典型的键-值存储区，由“非 SQL(noSQL)”操作进行填充。写入，也称 *put*，修改与某个给定键相关联的值。读取，也叫 *get*，返回某个键所对应的值。但这些一致性保证也可应用于采用其他读写操作类型的其他类别的复制数据存储系统，如文件系统和关系型数据库。这就是为什么这些一致性保证由写入（而不是数据值）进行定义的原因。例如，某个系统提供了增量操作或追加操作，那么，对对象进行的所有写入都可能对该对象的观测值构成影响，而不仅仅是最新的写入。此外，尽管本文的示例并不要求原子更新，这些保证也适用于那些访问多个对象的原子事务。

表 2 显示了各种一致性保证的典型性能和可用性。其中按照从劣到优分别对三个属性进行评级。一致性评级依据的是先前定义的一致性保证的强度。性能是指完成一个读操作所花的时间，即读取等待时间。可用性是指在服务器发生故障的情况下，一个读操作成功返回适当一致性数据的可能性。

从一致性的角度来看，强一致性是可取的，但其所提供的性能和可用性却是最差的，因为它通常需要读取指定主站或大部分副本。另一方面，最终一致性允许客户端从任意副本读取数据，但其提供的数据一致性最弱。性能与一致性之间呈负相关，这并不奇怪，因为较弱的一致性措施通常允许读请求被发送到更大范围的服务器。在有多个的数据足够新的服务器可供选择的情况下，客户端更有机会选中一台邻近的服务器。访问本地服务器与远程服务器的等待时间可能相差 100 倍。同理，较大的服务器选择余地，意味着客户端更容易找到一个（或通过定额配备一个）可达的服务器，从而获得更高的可用性。

每种一致性保证措施都提供了一致性、性能和可用性的一个独特组合。表 2 中每个单元的标签不是精确的科学（我可以就这个话题写上一整篇文章）。有人可能会说，标为“可”的一些项目实际上应该为“良”，反之亦然。的确，这些特征在某种程度上要取决于实现、部署和操作方面的细节。对于某些客户端，最终一致性读经常可能返回强一致性的结果，而且可能并不比强一致性读效率更高。^{3, 13} 但各种一致性保证措施之间的一般性比较，从定性上来看还是准确的。表 2 给出的信息至少说明，在选择采用特定一致性模型的特定数据复制方案时，人们需要进行仔细权衡。

表 2. 一致性、性能和价值权衡。

保证	一致性	性能	可用性
强一致性	优	劣	劣
最终一致性	劣	优	优
一致性前缀	可	良	优
限定过期法	良	可	劣
单调读	可	良	良
读自身写	可	可	可

图 1. 一场简化的棒球赛。

```
Write ( "visitors" , 0);
Write ( "home" , 0);
for inning = 1 ..9
  outs = 0;
  while outs < 3
    visiting player bats;
    for each run scored
      score = Read ( "visitors" );
      Write ( "visitors" , score + 1);
  outs = 0;
  while outs < 3
    home player bats;
    for each run scored
      score = Read ( "home" );
      Write ( "home" , score + 1);
end game;
```

虽然未提供任何证据，但我断言所有这些一致性保证均可成为同一存储系统的选项。实际上，我和同事已经在 MSR 硅谷实验室创建了此类系统的一个原型（但这是另一篇文章要讨论的话题）。在我们的这个系统中，要求不同一致性保证措施的客户端在执行读操作时，会有不同的性能和可用性体验，即使在访问共享数据时也是如此。让我们假定存在一个这样的存储系统，可以让客户端选择这六种读保证措施。下面我继续演示如何在使用这些措施 ... 以棒球为例。

以棒球赛为应用样例

对于不熟悉棒球，但喜欢读代码的读者，图 1 说明九局棒球赛的基本知识。比赛开始比分为 0-0。客队首先击球，然后连续击球，直到他们一方三人出局为止。随后主队击球，直到三人出局。比赛共进行九局。当然，这里省去了对于像我一样狂热的棒球爱好者来说非常重要的很多细节。不过，

图 2. 一个比赛示例的写序列。

```
Write ( "home" , 1)
Write ( "visitors" , 1)
Write ( "home" , 2)
Write ( "home" , 3)
Write ( "visitors" , 2)
Write ( "home" , 4)
Write ( "home" , 5)
```

对于本文来说，所有必要的东西的确都已解释清楚了。

假设比赛得分记录在两个对象中的一个键-值存储区，一个是“客队”得分的跑垒数，一个是“主队”得分的跑垒数。当一个队跑垒得一分时，读操作读取其当前得分，将返回值加 1，再将新值写回键-值存储区。

作为一个具体例子，请参看比赛示例的写入日志，如图 2 所示。在这场比赛中，主队先得一分，然后客队将比分扳平，接着主队连得两分，这样依次进行下去。

该序列写入的值可能来自棒球赛单行得分结果，各局比分如图

3 所示。这场假设的比赛现处于第七局中间（众所周知的第七局暂停），主队以 2-5 比分领先。

假设存有客队和主队总得分的键-值存储区驻留在云中，并在若干服务器间复制。不同的读保证措施可能导致客户端在比赛进行当中读到不同的得分。表 3 列出了读取客队和主队分数时，六种一致性保证措施中每一种措施可能返回的完整得分集合。注意：客队分数排在前面，所有可能出现的不同返回值用逗号隔开。

强一致性读只能返回一个结果，即当前比分，而最终一致性读可能返回 18 种可能比分中的一种。注意：由一对最终一致性读操作返回的很多比分结果，事实上从来没有在比赛中出现过。例如，读取客队得分时，可能会返回 2，而读取主队得分时，可能会返回 0，尽管主队在比赛中从未落后过。一致性前缀属性将比分结果限制为某段时间内真实存在过的比分。限定过期读操作所能返回的结果，明显取决于所期望的限定值。表 3 给出了限定为一局时的可能比分，也就是说，比分结果最多晚一局；当限定为七

图 3. 该比赛示例的单行得分。

	1	2	3	4	5	6	7	8	9	跑垒数
客队	0	0	1	0	1	0	0			2
主队	1	0	1	1	0	2				5

表 3. 对应每种一致性保证的可能读取的所有比分。

强一致性	2-5
最终一致性	0-0, 0-1, 0-2, 0-3, 0-4, 0-5, 1-0, 1-1, 1-2, 1-3, 1-4, 1-5, 2-0, 2-1, 2-2, 2-3, 2-4, 2-5
一致性前缀	0-0, 0-1, 1-1, 1-2, 1-3, 2-3, 2-4, 2-5
限定过期法	最多迟一局的比分：2-3, 2-4, 2-5
单调读	在读取 1-3 之后：1-3, 1-4, 1-5, 2-3, 2-4, 2-5
读自身写	对于数据写入者：2-5 对于数据写入者以外的任何人：0-0, 0-1, 0-2, 0-3, 0-4, 0-5, 1-0, 1-1, 1-2, 1-3, 1-4, 1-5, 2-0, 2-1, 2-2, 2-3, 2-4, 2-5

图 4. 记分员的角色。

```
score = Read ( "visitors" );
Write ( "visitors", score + 1);
```

局或更多局时，在本例中，其结果集与最终一致性的结果集相同。实际上，一个系统不太可能以“局”为单位来表示过期时间。因此，本例假设数据读取者请求 15 分钟的时间限定，而上一局正好持续了这么长时间。对于单调读，可能的返回值取决于过去读过哪些数据。对于读自身写，则取决于谁正在向键-值存储区写入数据；在本例中，假设所有的写操作均由单独一个客户端完成。

参与者的读需求

现在，让我们看看想知道一场棒球赛比分的各种人对数据一致性的需求。当然，每个人都可以通过执行强一致性读，来检索客队和主队的得分。正如我们在上一节所指出的那样，在这种情况下，只会有一种可能值被返回，即：当前比分。但如表 2 所示，要求强一致性的数据读取者，可能会经历更长的响应时间，甚至可能发现他们所请求的数据，因服务器临时故障或网络中断而变得不可用。本节的宗旨是评估每个参与者所需的最低一致性需求。这些客户端通过请求比强一致性弱的读保障措施，有可能体验到更好的性能和更高的数据可用性。

官方记分员。 官方记分员负责通过将比分写入持久性的键-值存储区，来维护比分。图 4 说明记分员在每次客队得分时采取的步骤；记分员在主队得分时采取的行动与此类似。注意：此代码为图 1 所示整个棒球赛代码的一个片段。

记分员的读操作需要什么一致性措施呢？毫无疑问，记分员需

图 5. 裁判的角色。

```
if first half of 9th inning complete then
  vScore = Read ( "visitors" );
  hScore = Read ( "home" );
  if vScore < hScore
    end game;
```

要读取最近的一次得分，然后将得分加 1，生成新的得分。否则，记分员就会冒险写入一个不正确得分，破坏整场比赛，还有可能激起愤怒球迷骚乱。假设主队此前已得五分，而且刚刚拿下第六分。如果采用最终一致性读，可能会返回从零到五的任何一个分值，如表 3 所示。也许，记分员非常幸运，读到了正确的得分，但这可不一定。

有趣的是，尽管记分员需要强一致性的数据，但他并不需要进行强一致性的读操作。由于记分员是唯一更新比赛得分的人，他可以请求读自身写保障措施，从而获得与强一致性读同样的效果。从本质上讲，记分员利用特定应用的知识，来获得较弱的一致性读所带来的好处，而不必真地放弃一致性。

这种差别似乎非常小，但实际上在实践中可能非常重要。存储系统在处理强一致性读的过程中，必须悲观地假定，来自世界任何地方的客户端都可能刚刚更新过数据。因此，系统必须访问大部分服务器（或一组固定的服务器），以确保客户端提交的读操作访问的是最新写入的数据。另一方面，系统在提供读自身写保证的过程中，只需记录客户端先前执行的写操作集合，并找到已完成所有这些写操作的服务器。¹¹ 在棒球赛中，之前得到的分数，也就是之前由记分员完成的写操作，可能在几分钟，甚至是几小时之前就已经发生。在这种情况下，几乎任何服务器都已收到之前的写入操作，而且能够应答下一个要求读自身写保证的读请求。

图 6. 电台体育记者的角色。

```
do {
  vScore = Read ( "visitors" );
  hScore = Read ( "home" );
  report vScore and hScore;
  sleep (30 minutes);
}
```

裁判。 裁判是指在本垒板后主持棒球赛的人。裁判在大多数情况下并不真正关心比赛的当前比分。只有一个例外，就是在第 9 局上半场之后，也就是客队已击完球，主队正准备击球的那个时刻。由于这是最后一局（每个队都不能得负分），主队如果比分领先的话，就表明已经赢了；因此，在某些比赛中，主队可以跳过最后一局的击球。需要作此决定的裁判所对应的代码如图 5 所示。

在第 9 局进行中，裁判要读取得分时，的确需要读取当前得分。否则，如果他错以为主队领先，就可能会过早地结束比赛，或者让主队进行不必要的击球。与记分员不同，裁判从来不会写得分；他只是读取由官方记分员写入的值。因此，为了获得最新信息，裁判必须采用强一致性读。

电台记者。 在美国的大部分地区，电台会定期宣布正在进行的或已结束的比赛的比分。例如，旧金山地区的 KCBS 电台每隔 30 分钟就会播报一次体育新闻。电台记者完成图 6 所示的步骤。一个类似的，也许更为现代的例子是，当观众正在观看 ESPN 时，电视屏幕底部滚动显示体育比赛成绩。

如果电台记者播报的不是最新的比分，也没有问题。人们都已习惯收听旧闻。因此，对于他的

图 7. 体育新闻记者的角色。

```
While not end of game {
    drink beer;
    smoke cigar;
}
go out to dinner;
vScore = Read ( "visitors" );
hScore = Read ( "home" );
write article;
```

图 8. 统计员的角色。

```
Wait for end of game;
score = Read ( "home" );
stat = Read ( "season-runs" );
Write ( "season-runs", stat +
score);
```

图 9. 统计监察员的角色。

```
do {
    stat = Read ( "season-runs" );
    discuss stats with friends;
    sleep (1 day);
}
```

读操作，某种形式的最终一致性也是可以的。哪种一致性保证是电台记者所希望的呢？

如表 3 所示，采用最弱保证的读，即最终一致性的读，可能会返回从来不存在的比分。对于图 3 给出的单行得分示例，此种读操作可能会返回一个客队 1-0 领先的比分，尽管客队实际上从未领先过。电台记者不希望播报这种虚假的比分。因此，记者希望他两次都从快照中读取数据，快照中包含由记分员写入的值的一致性前缀。这使得记者能够读到之前某个时间的实际比分，但不必是当前比分。

然而，仅读一致性前缀还不够。对于图 3 所示的单行得分，

记者可能读到一个 2-5 的比分，即当前比分，等过了 30 分钟，又读到一个 1-3 的比分。例如，如果记者碰巧从主服务器上读取数据，而后又从可能位于远程数据中心的另一台服务器上读取数据，而该服务器已从主站上断开，尚未接收最新的写入数据，那么就可能出现上述情况。由于大家都知道棒球比分是单调递增的，如果在随后的新闻报道中依次播出 2-5 和 1-3 的比分，会让人们觉得记者非常愚笨。如果记者除了要求一致性前缀之外，还要求单调读保证，那么就可以避免出现这种情形。请注意：两种保证措施任中一种都不能独立满足要求。

另外，记者请求限定期限不超过 30 分钟的限定过期读，配合单调读，可以得到同样的效果。这将确保记者观察到的比分最多过时 30 分钟。由于记者每 30 分钟才读取一次数据，他接收的比分一定越来越新。当然，记者可以要求更紧的限定，比如五分钟，以便获得具有合理及时性的比分。

体育新闻记者。另一个有趣的人物就是体育新闻记者。他在观看比赛后，要写一篇报导发表在晨报上或上传到某个网站。不同体育新闻记者的表现可能也不相同，但通过观察（我也当过体育新闻记者），我发现他们的行为通常如图 7 所示。

体育新闻记者可能并不急着写他的文章。在这个例子中，他悠闲地外出吃晚饭，然后坐下来对比赛进行总结。他肯定要保证

他报导的比分是比赛的最终正确得分。因此，他想得到强一致性读的效果。不过，他并不需要支付这笔开销。如果体育新闻记者知道他在比赛结束后花了一个小时吃晚饭，那么他也一定知道，距离记分员最后更新比分至少也有一个小时了。因此，为限定过期读设置一小时的限定，足以保证体育新闻记者能够读到最终比分。在实际应用中，任何服务器都应该能回应此类读操作。实际上，最终一致性读很可能在一个小时后返回正确的比分，但是唯一能够让体育新闻记者 100% 确定自己得到最终比分的方法，就是请求采用限定过期法。

统计员。球队统计员负责跟踪球队和每个球员各赛季的统计数据。例如，统计员可能计算本赛季她的球队的得分总数。假设这些统计数据也保存在持久性的键-值存储区。如图 8 所示，主队统计员在每场比赛结束后的某个时间，将得分加到上个赛季得分总数上，然后将这个新值写入数据存储区。

当统计员读取当天的球队得分时，她需要确信自己得到的是最终比分。因此，她需要执行一个强一致性读操作。如果统计员在比赛结束后等待一段时间，那么，限定过期读也可能实现同样的效果（正如前面针对体育新闻记者的讨论）。

当读取本赛季当前的统计数据时，也就是图 8 所示的第二个读操作，统计员也希望采用强一致性读操作。如果返回的是旧统计数字，那么，写回的更新值会令球队的总得分减少。由于统计员是唯一将统计数据写入数据存储区的人，因此，她可以采用读自身写保证，以获得最新值（如前面讨论的那样）。

统计监察员。定期检查球队赛季统计数据的其他人员通常会满

表 4 棒球赛参与者的读保证方式。

官方记分员	读自身写
裁判	强一致性
电台记者	一致性前缀与单调读
体育新闻记者	限定过期法
统计员	强一致性、读自身写
统计监察员	最终一致性

足于最终一致性。统计数据每天仅更新一次，略显过时的数字还可以接受。例如，一个球迷询问本赛季他的球队的总得分，就可以通过最终一致性读，得到一个合理的答案，如图9所示。

结论

显然，存储棒球赛得分并非云存储系统的杀手级应用。对于从一个简单例子得出结论，我们应当谨慎。但也许我们可以从中得到一些教训。

对于上一节讨论过的各类棒球赛参与人员所希望采取的一致性保证，表4进行了总结。记住：所列举的一致性措施并非唯一可接受的措施。具体而言，每个参与者采用强一致性是没有问题的，但是，放松对一致性读的要求，很可能会获得更好的性能和更高的可用性。此外，存储系统可以更好地做到跨服务器读工作负荷均衡，这是因为，在选择用于应答弱一致性读请求的服务器方面有了更大的灵活性。

这些参与者可以看作是不同的应用程序，它们访问共享的数据：棒球赛比分。对于某些情形，如记分员和体育新闻记者，基于特定应用知识，数据读取者很清楚：即使他采用读自身写或限定过期保证方式发出弱一致性读请求，他也可以得到强一致性数据。对于另一些情形，如电台记者，多种保证方式必须组合起来使用，才能满足数据读取者的需求。对于其他情形，如统计员，希望采用不同的保证方式读取不同的数据对象。

我从本练习中得出了四个主要结论：

► **本文提出的所有六种一致性保证都非常有用。** 请注意：每种保证在表4中至少出现一次。除了一类客户端外，仅提供最终一致性的系统将无法满足其他所有类型客户端的需求；除了两类客户端

外，仅提供强一致性的系统可能会在其他所有情形下表现欠佳。

► **不同的客户端在访问相同的数据时，可能希望采用不同的一致性措施。** 系统经常会将一个具体的一致性绑定到一个特定的数据集或数据类。例如，人们普遍认为，银行数据必须是强一致性的，而购物车数据仅需要最终一致性即可。棒球赛的例子表明，人们所期望的一致性既决于数据类型，也取决于谁在读取数据。

► **即使是简单的数据库，各种用户也可能对一致性有不同的需求。** 棒球赛比分是可以想象的最简单的一种数据库，仅由两个数字组成。但这有效地阐释了不同一致性选项的价值所在。

► **客户端应能选择自己想要的一致性。** 系统可能无法预测或确定指定应用程序或客户端所需的一致性。首选的一致性选项往往取决于数据的使用方式。此外，在读最新的数据时，了解谁写入的数据，或者知道数据最后写入的时间，有时可能允许客户端执行松弛一致性读，从而获得相应的好处。

反对提供最终一致性方式的人们所持的主要论点经常是，它增加了应用程序开发人员的负担。虽然这可能是事实，但不必有过多的额外负担。第一步是定义开发人员可以理解的一致性保证；请注意，表1中的六种保证分别是用寥寥数语加以描述的。让存储系统依照严格顺序执行写操作，就可以使应用程序开发人员避免对并发写入所造成的更新冲突问题进行处理时所遇到的复杂性。这样开发人员的工作就只剩下选择自己想要的读一致性措施。要做出这种选择，就需要他们对其应用程序的语义有深入了解，但不必改变程序的基本结构。上一节提供的几个代码片段，都不需要增加额外的代码行来专门处理过时的数据。

仅提供强一致性的云存储系统，让开发人员能够很容易编写出

正确的程序，但可能会错失松弛一致性所带来的好处。一致性、性能和可用性之间固有的权衡是实际存在的，随着复制服务地理范围的扩大，问题有可能变得更加突出。这表明云存储系统至少应考虑提供更多的读一致性选择。有些云服务提供商已经提供强一致性和最终一致性两种读操作，但从文中可以看出，他们的最终一致性模型对应用程序来说可能并不理想。云存储客户端通过选择几种一致性保证，并从不同的副本读取数据，不仅可以让各种各样的应用程序受益，而且可以获得更高的资源利用率，并节约成本。

参考资料

1. Abadi, D. Consistency tradeoffs in modern distributed database system design. *IEEE Computer*, (Feb. 2012).
2. Amazon. Amazon DynamoDB; <http://aws.amazon.com/dynamodb/>.
3. Anderson, E., Li, X., Shah, M., Tucek, J. and Wylie, J. What consistency does your key-value store actually provide? In *Proceedings of the Usenix Workshop on Hot Topics in Systems Dependability*, (2010).
4. Bailis, P., Venkataraman, S., Franklin, M., Hellerstein, J. and Stoica, I. Probabilistically bounded staleness for practical partial quorums. In *Proceedings VLDB Endowment*, (Aug. 2012).
5. Brewer, E. CAP twelve years later: How the "rules" have changed. *IEEE Computer*, (Feb. 2012).
6. Calder, B. et. al. Windows azure storage: A highly available cloud storage service with strong consistency. In *Proceedings ACM Symposium on Operating Systems Principles*, (Oct. 2011).
7. Cooper, B., Ramakrishnan, R., Srivastava, U., Silberstein, A., Bohannon, P., Jacobsen, H.A., Puz, N., Weaver, D. and Yerneni, R. PNUTS: Yahoo!'s hosted data serving platform. In *Proceedings International Conference on Very Large Data Bases*, (Aug. 2008).
8. Google. Read Consistency & Deadlines: More Control of Your Datastore. Google App Engine Blog, Mar. 2010; <http://googleappengine.blogspot.com/2010/03/read-consistency-deadlines-more-control.html>.
9. Kraska, T., Hentschel, M., Alonso, G. and Kossmann, D. Consistency rationing in the cloud: Pay only when it matters. In *Proceedings International Conference on Very Large Data Bases*, (Aug. 2009).
10. Saito, Y. and Shapiro, M. Optimistic replication. *ACM Computing Surveys*, (Mar. 2005).
11. Terry, D., Demers, A., Petersen, K., Spreitzer, M., Theimer, M., and Welch, B. Session guarantees for weakly consistent replicated data. In *Proceedings IEEE International Conference on Parallel and Distributed Information Systems*, (1994).
12. Vogels, W. Eventually consistent. *Commun. ACM*, (Jan. 2009).
13. Wada, H., Fekete, A., Zhao, L., Lee, K. and Liu, A. Data consistency properties and the trade-offs in commercial cloud storages: The consumers' perspective. In *Proceedings CIDR*, (Jan. 2011).

Doug Terry (terry@microsoft.com) 是位于加州山景城 (Mountain View) 的微软研究院硅谷实验室首席研究员。

版权所有 / 作者所有。发表权已授予 ACM。\$15.00

A broader class of consistency guarantees can, and perhaps should, be offered to clients that read shared data.

BY DOUG TERRY

Replicated Data Consistency Explained Through Baseball

REPLICATED STORAGE SYSTEMS for the cloud deliver different consistency guarantees to applications that are reading data. Invariably, cloud storage providers redundantly store data on multiple machines so that data remains available in the face of unavoidable failures. Replicating data across datacenters is not uncommon, allowing the data to survive complete site outages. However, the replicas are not always kept perfectly synchronized. Thus, clients that read the same data object from different servers can potentially receive different versions.

Some systems, like Microsoft's Windows Azure, provide only strongly consistent storage services to

their applications.⁵ These services ensure clients of Windows Azure Storage always see the latest value that was written for a data object. While strong consistency is desirable and reasonable to provide within a datacenter, it raises concerns as systems start to offer geo-replicated services that span multiple datacenters on multiple continents.

Many cloud storage systems, such as the Amazon Simple Storage Service (S3), were designed with weak consistency based on the belief that strong consistency is too expensive in large systems. The designers chose to relax consistency in order to obtain better performance and availability. In such systems, clients may perform read operations that return stale data. The data returned by a read operation is the value of the object at *some past point in time* but not necessarily the latest value. This occurs, for instance, when the read operation is directed to a replica that has not yet received all of the writes that were accepted by some other replica. Such systems are said to be *eventually consistent*.¹²

Recent systems, recognizing the need to support different classes of applications, have been designed with a choice of operations for accessing cloud storage. Amazon's DynamoDB, for example, provides both *eventually consistent reads* and *strongly consistent reads*, with the latter experiencing a higher read latency and a twofold reduction in read throughput.¹ Amazon SimpleDB offers the same choices for clients that

» key insights

- Although replicated cloud services generally offer strong or eventual consistency, intermediate consistency guarantees may better meet an application's needs.
- Consistency guarantees can be defined in an implementation-independent manner and chosen for each read operation.
- Dealing with relaxed consistency need not place an excessive burden on application developers.



read data. Similarly, the Google App Engine Datastore added eventually consistent reads to complement its default strong consistency.⁸ PNUTS, which underlies many of Yahoo's Web services, provides three types of read operations: *read-any*, *read-critical*, and *read-latest*.⁷ Modern quorum-based storage systems allow clients to choose between strong and eventual consistency by selecting different read and write quorums.⁴

In the research community over the past 30 years, a number of consistency models have been proposed for distributed and replicated systems.¹⁰ These offer consistency guarantees that lie somewhere in between strong consistency and eventual consistency. For example, a system might guarantee that a client sees data that is no more than five minutes out of date or that a client always observes the results of its own writes. Actually, some consistency

models are even weaker than eventual consistency, but those I ignore as being less than useful.

The reason for exploring different consistency models is that there are fundamental trade-offs between consistency, performance, and availability.^{9,10,12,13} Offering stronger consistency generally results in lower performance and reduced availability for reads or writes or both. The CAP theorem has proven that, for systems

Table 1. Six consistency guarantees.

Strong Consistency	See all previous writes.
Eventual Consistency	See subset of previous writes.
Consistent Prefix	See initial sequence of writes.
Bounded Staleness	See all “old” writes.
Monotonic Reads	See increasing subset of writes.
Read My Writes	See all writes performed by reader.

that must tolerate network partitions, designers must choose between consistency and availability.⁵ In practice, latency is an equally important consideration.⁴ Each proposed consistency model occupies some point in the complex space of trade-offs.

Are different consistencies useful in practice? Can application developers cope with eventual consistency? Should cloud storage systems offer an even greater choice of consistency than the consistent and eventually consistent reads offered by some of today’s services?

This article attempts to answer these questions, at least partially, by examining an example (but clearly fictitious) application: the game of baseball. In particular, I explore the needs of different people who access the score of a baseball game, including the scorekeeper, umpire, radio reporter, sportswriter, and statistician. Supposing the score is stored in a cloud-based, replicated storage service, I show eventual consistency is insufficient for most of the participants, but strong consistency is not needed either. Most participants benefit from some intermediate consistency guarantee.

The next section defines six possible consistency guarantees for read operations. Then I present an algorithm that emulates a baseball game, indicating where data is written and read, and I enumerate the results that might be returned when reading the score with different guarantees. I also examine the roles of various people who want to access the baseball score and the read consistency that each desires and draw conclusions from this simple example.

Read Consistency Guarantees

While replicated systems have provided many types of data consistency over

the past 30 years, and a wide variety of consistency models have been explored in the computer science research community, many of these are tied to specific implementations. Frequently, one needs to understand how a system operates in order to understand what consistency it provides in what situations. This places an unfortunate burden on those who develop applications on top of such storage systems.

The six consistency guarantees I advocate here can be described in a simple, implementation-independent way. This not only benefits application developers but also can permit flexibility in the design, operation, and evolution of the underlying storage system.

These consistency guarantees are based on a simple model in which clients perform *read* and *write* operations to a data store. Multiple clients may concurrently access shared information, such as social network graphs, news feeds, photos, shopping carts, or financial records. The data is replicated among a set of servers, but the details of the replication protocol are hidden from clients. A write is any operation that updates one or more data objects. Writes are eventually received at all servers and performed in the same order. This order is consistent with the order in which clients submit write operations. In practice, the order could be enforced by performing all writes at a master server or by having servers run a consensus protocol to reach agreement on the global order. Reads return the values of one or more data objects that were previously written, though not necessarily the latest values. Each read operation can request a consistency guarantee, which dictates the set of allowable return values. Each guarantee is defined by the set of previous

writes whose results are visible to a read operation. Table 1 summarizes these six consistency guarantees.

Strong consistency is particularly easy to understand. It guarantees a read operation returns the value that was last written for a given object. If write operations can modify or extend portions of a data object, such as appending data to a log, then the read returns the result of applying all writes to that object. In other words, a read observes the effects of *all* previously completed writes.

Eventual consistency is the weakest of the guarantees, meaning it allows the greatest set of possible return values. For whole-object writes, an eventually consistent read can return any value for a data object that was written in the past. More generally, such a read can return results from a replica that has received an arbitrary subset of the writes to the data object being read. The term “eventual” consistency derives from the fact that each replica eventually receives each write operation, and if clients stopped performing writes then read operations would eventually return an object’s latest value.

By requesting a *consistent prefix*, a reader is guaranteed to observe an ordered sequence of writes starting with the first write to a data store. For example, the read may be answered by a replica that receives writes in order from a master replica but has not yet received some recent writes. In other words, the reader sees a version of the data store that existed at the master at some time in the past. This is similar to the “snapshot isolation” consistency offered by many database management systems. For reads to a single data object in a system where write operations completely overwrite previous values of an object, even eventual consistency reads observe a consistent prefix. The main benefit of requesting a consistent prefix arises when reading multiple data objects or when write operations incrementally update an object.

Bounded staleness ensures read results are not too out of date. Typically, staleness is defined by a time period T , say five minutes. The storage system guarantees a read operation will return any values written more than T minutes ago or more recently writ-

ten values. Alternative, some systems have defined staleness in terms of the number of missing writes or even the amount of inaccuracy in a data value. I find that time-bounded staleness is the most natural concept for application developers.

Monotonic reads is a property that applies to a sequence of read operations performed by a given storage system client. As such, it is called a “session guarantee.”¹¹ With monotonic reads, a client can read arbitrarily stale data, as with eventual consistency, but is guaranteed to observe a data store that is increasingly up to date over time. In particular, if the client issues a read operation and then later issues another read to the same object(s), the second read will return the same value(s) or a more recently written value.

Read my writes is a property that also applies to a sequence of operations performed by a single client. It guarantees the effects of all writes that were performed by the client are visible to the client’s subsequent reads. If a client writes a new value for a data object and then reads this object, the read will return the value that was last written by the client (or some other value that was later written by a different client). For clients that have issued no writes, the guarantee is the same as eventual consistency. (Note: In previous articles this has been called “Read Your Writes,”¹¹ but I have chosen to rename it to more accurately describe the guarantee from the client’s viewpoint.)

These last four read guarantees are all a form of eventual consistency but stronger than the eventual consistency model that is typically provided in cloud storage systems. The “strength” of a consistency guarantee does not depend on when and how writes propagate between servers, but rather is defined by the size of the set of allowable results for a read operation. Smaller sets of possible read results indicate stronger consistency. When requesting strong consistency, there is a single value that must be returned, the latest value that was written. For an object that has been updated many times, an eventually consistent read can return one of many suitable values. Of the four intermediate guarantees, none is stron-

ger than any of the others, meaning each might have a different set of possible responses to a read operation. In some cases, as will be shown later, applications may want to request multiple guarantees. For example, a client could request both monotonic reads and read my writes so that it observes a data store that is consistent with its own actions.¹¹

In this article, the data store used for baseball scores is a traditional key-value store, popularized by the “noSQL” movement. Writes, also called *puts*, modify the value associated with a given key. Reads, also called *gets*, return the value for a key. However, these guarantees can apply to other types of replicated data stores with other types of read and write operations, such as file systems and relational databases. This is why the guarantees are defined in terms of writes rather than data values. For example, in a system that offers an increment or an append operation, all writes performed on an object contribute to the object’s observed value, not just the latest write. Moreover, the guarantees could apply to atomic transactions that access multiple objects, though the examples in this article do not require atomic updates.

Table 2 shows the performance and availability typically associated with each consistency guarantee. It rates the three properties on a scale from poor to excellent. Consistency ratings are based on the strength of the consistency guarantee as previously defined. Performance refers to the time it takes to complete a read operation, that is, the read latency. Availability is the likelihood of a read operation successfully returning suitably consistent data in the presence of server failures.

Strong consistency is desirable

from a consistency viewpoint but offers the worst performance and availability since it generally requires reading from a designated primary site or from a majority of replicas. Eventual consistency, on the other hand, allows clients to read from any replica, but offers the weakest consistency. The inverse correlation between performance and consistency is not surprising since weaker forms of consistency generally permit read requests to be sent to a wider set of servers. With more choices of servers that are sufficiently up to date, clients are more able to choose a nearby server. The latency difference between accessing a local rather than a remote server can be a factor of 100. Similarly, a larger choice of servers means a client is more likely to find one (or a quorum) that is reachable, resulting in higher availability.

Each guarantee offers a unique combination of consistency, performance, and availability. Labeling each cell in Table 2 is not an exact science (and I could devote a whole article to this topic). One might argue that some entry listed as “okay” should really be “good”, or vice versa, and indeed the characteristics do depend to some extent on implementation, deployment, and operating details. For some clients, eventually consistent reads may often return strongly consistent results, and may not be any more efficient than strongly consistent reads.^{3,13} But, the general comparisons between the various consistency guarantees are qualitatively accurate. The bottom line is that one faces substantial trade-offs when choosing a particular replication scheme with a particular consistency model.

Without offering any evidence, I assert that all of these guarantees can be provided as choices within

Table 2. Consistency, performance, and availability trade-offs.

Guarantee	Consistency	Performance	Availability
Strong Consistency	excellent	poor	poor
Eventual Consistency	poor	excellent	excellent
Consistent Prefix	okay	good	excellent
Bounded Staleness	good	okay	poor
Monotonic Reads	okay	good	good
Read My Writes	okay	okay	okay

Figure 1. A simplified baseball game.

```

Write ("visitors", 0);
Write ("home", 0);
for inning = 1 .. 9
  outs = 0;
  while outs < 3
    visiting player bats;
    for each run scored
      score = Read ("visitors");
      Write ("visitors", score + 1);
  outs = 0;
  while outs < 3
    home player bats;
    for each run scored
      score = Read ("home");
      Write ("home", score + 1);
end game;

```

the same storage system. In fact, my colleagues and I at the MSR Silicon Valley Lab have built a prototype of such a system (but that is the topic for another article). In our system, clients requesting different consistency guarantees experience different performance and availability for the read operations they perform, even when accessing shared data. Here, let's assume the existence of a storage system that offers its clients a choice of these six read guarantees. I proceed to show how they would be used...in baseball.

Baseball as a Sample Application

For those readers who are not familiar with baseball, but who love to read code, Figure 1 illustrates the basics of a nine-inning baseball game. The game starts with the score of 0-0. The visitors bat first and remain at bat until they make three outs. Then the home team bats until it makes three outs. This continues for nine innings. Granted, this leaves out many of the subtleties that are dear to baseball aficionados, like myself. But it does explain all that is needed for this article.

Assume the score of the game is recorded in a key-value store in two

Figure 2. Sequence of writes for a sample game.

```

Write ("home", 1)
Write ("visitors", 1)
Write ("home", 2)
Write ("home", 3)
Write ("visitors", 2)
Write ("home", 4)
Write ("home", 5)

```

objects, one for the number of runs scored by the "visitors" and one for the "home" team's runs. When a team scores a run, a read operation is performed on its current score, the returned value is incremented by one, and the new value is written back to the key-value store.

As a concrete example, consider the write log for a sample game as shown in Figure 2. In this game, the home team scored first, then the visitors tied the game, then the home team scored twice more, and so on.

This sequence of writes could be from a baseball game with the inning-by-inning line score that is illustrated in Figure 3. This hypothetical game is currently in the middle of the seventh inning (the proverbial seventh-inning

stretch), and the home team is winning 2-5.

Suppose the key-value store that holds the visitors and home team's run totals resides in the cloud and is replicated among a number of servers. Different read guarantees may result in clients reading different scores for this game that is in progress. Table 3 lists the complete set of scores that could be returned by reading the visitors and home scores with each of the six consistency guarantees. Note that the visitors' score is listed first, and different possible return values are separated by comas.

A strong consistency read can only return one result, the current score, whereas an eventual consistency read can return one of 18 possible scores. Observe that many of the scores that can be returned by a pair of eventually consistent reads are ones that were never the actual score. For example, reading the visitors' score may return two and reading the home team's score may return zero, even though the home team never trailed. The consistent prefix property limits the result to scores that actually existed at some time. The results that can be returned by a bounded staleness read clearly depend on the desired bound. Table 3 illustrates the possible scores for a bound of one inning out of date; for a bound of seven innings or more,

Figure 3. The line score for this sample game.

	1	2	3	4	5	6	7	8	9	RUNS
Visitors	0	0	1	0	1	0	0			2
Home	1	0	1	1	0	2				5

Table 3. Possible scores read for each consistency guarantee.

Strong Consistency	2-5
Eventual Consistency	0-0, 0-1, 0-2, 0-3, 0-4, 0-5, 1-0, 1-1, 1-2, 1-3, 1-4, 1-5, 2-0, 2-1, 2-2, 2-3, 2-4, 2-5
Consistent Prefix	0-0, 0-1, 1-1, 1-2, 1-3, 2-3, 2-4, 2-5
Bounded Staleness	scores that are at most one inning out-of-date: 2-3, 2-4, 2-5
Monotonic Reads	after reading 1-3: 1-3, 1-4, 1-5, 2-3, 2-4, 2-5
Read My Writes	for the writer: 2-5 for anyone other than the writer: 0-0, 0-1, 0-2, 0-3, 0-4, 0-5, 1-0, 1-1, 1-2, 1-3, 1-4, 1-5, 2-0, 2-1, 2-2, 2-3, 2-4, 2-5

Figure 4. Role of the scorekeeper.

```
score = Read ("visitors");
Write ("visitors", score + 1);
```

the result set is the same as for eventual consistency in this example. In practice, a system is unlikely to express staleness bounds in units of “innings.” So, for this example, assume the reader requested a bound of 15 minutes and the previous inning lasted exactly that long. For monotonic reads, the possible return values depend on what has been read in the past. For read my writes they depend on who is writing to the key-value store; in this example, assume all of the writes were performed by a single client.

Read Requirements for Participants

Now, let’s examine the consistency needs of a variety of people involved in a baseball game who want to read the score. Certainly, each of these folks could perform a strongly consistent read to retrieve the visiting and home team’s score. In this case, as pointed out in the previous section, only one possible value would be returned: the current score. However, as shown in Table 2, readers requesting strong consistency will likely receive longer response times and may even find that the data they are requesting is not currently available due to temporary server failures or network outages. The point of this section is to evaluate, for each participant, the minimum consistency that is required. By requesting read guarantees that are weaker than strong consistency, these clients are likely to experience performance benefits and higher availability.

Official scorekeeper. The official scorekeeper is responsible for maintaining the score of the game by writing it to the persistent key-value store. Figure 4 illustrates the steps taken by the scorekeeper each time the visiting team scores a run; his action when the home team scores is similar. Note that this code is a snippet of the overall baseball game code that was presented in Figure 1.

What consistency does the scorekeeper require for his read operations? Undoubtedly, the scorekeeper

Figure 5. Role of the umpire.

```
if first half of 9th inning complete then
  vScore = Read ("visitors");
  hScore = Read ("home");
  if vScore < hScore
    end game;
```

needs to read the most up-to-date previous score before adding one to produce the new score. Otherwise, the scorekeeper runs the risk of writing an incorrect score and undermining the game, not to mention inciting a mob of angry baseball fans. Suppose the home team had previously scored five runs and just scored the sixth. Doing an eventual consistency read, as shown in Table 3, could return a score of anything from zero to five. Perhaps, the scorekeeper would get lucky and receive the correct score in response to his read, but he should not count on it.

Interestingly, while the scorekeeper requires strongly consistent data, he does not need to perform strong consistency reads. Since the scorekeeper is the only person who updates the score, he can request the read my writes guarantee and receive the same effect as a strong read. Essentially, the scorekeeper uses application-specific knowledge to obtain the benefits of a weaker consistency read without actually giving up any consistency.

This might seem like a subtle distinction, but, in fact, could be quite significant in practice. In processing a strong consistency read the storage system must pessimistically assume that some client, anywhere in the world, may have just updated the data. The system therefore must access a majority of servers (or a fixed set of servers) in order to ensure the most recently written data is accessed by the submitted read operation. In providing the read my writes guarantee, on the other hand, the system simply needs to record the set of writes that were previously performed by the client and find some server that has seen all of these writes.¹¹ In a baseball game, the previous run that was scored, and hence the previous write that was performed by the scorekeeper, may have happened many minutes or even hours ago. In this case, almost any server will have received the previous write and be able to answer the

Figure 6. Role of the radio sports reporter.

```
do {
  vScore = Read ("visitors");
  hScore = Read ("home");
  report vScore and hScore;
  sleep (30 minutes);
}
```

next read that requests the read my writes guarantee.

Umpire. The umpire is the person who officiates a baseball game from behind home plate. The umpire, for the most part, does not actually care about the current score of the game. The one exception comes after the top half of the 9th inning, that is, after the visiting team has batted and the home team is about to bat. Since this is the last inning (and a team cannot score negative runs), the home team has already won if they are ahead in the score; thus, the home team can and does skip its last at bat in some games. The code for the umpire who needs to make this determination is illustrated in Figure 5.

When accessing the score during the 9th inning, the umpire does need to read the current score. Otherwise, he might end the game early, if he incorrectly believes the home team to be ahead, or make the home team bat unnecessarily. Unlike the scorekeeper, the umpire never writes the score; he simply reads the values that were written by the official scorekeeper. Thus, in order to receive up-to-date information, the umpire must perform strong consistency reads.

Radio reporter. In most areas of the U.S., radio stations periodically announce the scores of games that are in progress or have completed. In the San Francisco area, for example, KCBS reports sports news every 30 minutes. The radio reporter performs the steps outlined in Figure 6. A similar, perhaps more modern, example is the sports scores that scroll across the bottom of the TV screen while viewers are watching ESPN.

Figure 7. Role of the sportswriter.

```

While not end of game {
    drink beer;
    smoke cigar;
}
go out to dinner;
vScore = Read ("visitors");
hScore = Read ("home");
write article;

```

Figure 8. Role of the statistician.

```

Wait for end of game;
score = Read ("home");
stat = Read ("season-runs");
Write ("season-runs", stat + score);

```

Figure 9. Role of the stat watcher.

```

do {
    stat = Read ("season-runs");
    discuss stats with friends;
    sleep (1 day);
}

```

If the radio reporter broadcasts scores that are not completely up to date, that is okay. People are accustomed to receiving old news. Thus, some form of eventual consistency is fine for the reads he performs. But what guarantees, if any, are desirable?

As shown in Table 3, the read with the weakest guarantee, an eventual consistency read, may return scores that never existed. For the sample line score given in Figure 3, such a read might return a score with the visitors leading 1-0, even though the visiting team has never actually been in the lead. The radio reporter does not want to report such fictitious scores. Thus, the reporter wants both his reads to be performed on a snapshot that hold a consistent prefix of the writes that were performed by the scorekeeper. This allows the reporter

to read the score that existed at some time, without necessarily reading the current score.

But reading a consistent prefix is not sufficient. For the line score in Figure 3, the reporter could read a score of 2-5, the current score, and then, 30 minutes later, read a score of 1-3. This might happen, for instance, if the reporter happens to read from a primary server and later reads from another server, perhaps in a remote datacenter, that has been disconnected from the primary and has yet to receive the latest writes. Since everyone knows that baseball scores are monotonically increasing, reporting scores of 2-5 and 1-3 in subsequent news reports would make the reporter look foolish. This can be avoided if the reporter requests the monotonic reads guarantee in addition to requesting a consistent prefix. Observe that neither guarantee is sufficient by itself.

Alternatively, the reporter could obtain the same effect as a monotonic read by requesting bounded staleness with a bound of less than 30 minutes. This would ensure the reporter observes scores that are at most 30 minutes out of date. Since the reporter only reads data every 30 minutes, he must receive scores that are increasingly up to date. Of course, the reporter could ask for a tighter bound, say five minutes, to get scores that are reasonably timely.

Sportswriter. Another interesting person is the sportswriter who watches the game and later writes an article that appears in the morning paper or that is posted on some website. Different sportswriters may behave differently, but my observations (from having been a sportswriter) is they often act as in Figure 7.

The sportswriter may be in no hurry to write his article. In this example, he

goes out to a leisurely dinner before sitting down to summarize the game. He certainly wants to make sure that he reports the correct final score for the game. So, he wants the effect of a strong consistency read. However, he does not need to pay the cost. If the sportswriter knows he spent an hour eating dinner after the game ended, then he also knows it has been at least an hour since the scorekeeper last updated the score. Thus, a bounded staleness read with a bound of one hour is sufficient to ensure the sportswriter reads the final score. In practice, any server should be able to answer such a read. In fact, an eventual consistency read is likely to return the correct score after an hour, but requesting bounded staleness is the only way for the sportswriter to be 100% certain he is obtaining the final score.

Statistician. The team statistician is responsible for keeping track of the season-long statistics for the team and for individual players. For example, the statistician might tally the total number of runs scored by her team this season. Suppose these statistics are also saved in the persistent key-value store. As shown in Figure 8, the home team's statistician, sometime after each game has ended, adds the runs scored to the previous season total and writes this new value back into the data store.

When reading the team's score from today, the statistician wants to be sure to obtain the final score. Thus, she needs to perform a strong consistency read. If the statistician waits for some time after the game, then a bounded staleness read may achieve the same effect (as discussed earlier for the sportswriter).

When reading the current statistics for the season, that is, for the second read operation in Figure 8, the statistician also wants strong consistency. If an old statistic is returned, then the updated value written back will undercount the team's total runs. Since the statistician is the only person who writes statistics into the data store, she can use the read my writes guarantee to get the latest value (as discussed previously).

Stat watcher. Others who periodically check on the team's season statistics are usually content with eventual consistency. The statistical data is only updated once per day, and numbers

Table 4. Read guarantees for baseball participants.

Official scorekeeper	Read My Writes
Umpire	Strong Consistency
Radio reporter	Consistent Prefix & Monotonic Reads
Sportswriter	Bounded Staleness
Statistician	Strong Consistency, Read My Writes
Stat watcher	Eventual Consistency

that are slightly out of date are okay. For example, a fan inquiring about the total number of runs scored by his team this season, as shown in Figure 9, can perform an eventual consistency read to get a reasonable answer.

Conclusion

Clearly, storing baseball scores is not the killer application for cloud storage systems. And we should be cautious about drawing conclusions from one simple example. But perhaps some lessons can be learned.

Table 4 summarizes the consistency guarantees desired by the variety of baseball participants that were discussed in the previous section. Recall that the listed consistencies are not the only acceptable ones. In particular, each participant would be okay with strong consistency, but, by relaxing the consistency requested for his reads, he will likely observe better performance and availability. Additionally, the storage system may be able to better balance the read workload across servers since it has more flexibility in selecting servers to answer weak consistency read requests.

These participants can be thought of as different applications that are accessing shared data: the baseball score. In some cases, such as for the scorekeeper and sportswriter, the reader, based on application-specific knowledge, knows he can obtain strongly consistent data even when issuing a weakly consistent read using a read my writes or bounded staleness guarantee. In some cases, such as the radio reporter, multiple guarantees must be combined to meet the reader's needs. In other cases, such as the statistician, different guarantees are desired for reads to different data objects.

I draw four main conclusions from this exercise:

► **All of the six presented consistency guarantees are useful.** Observe that each guarantee appears at least once in Table 4. Systems that offer only eventual consistency would fail to meet the needs of all but one of these clients, and systems that offer only strong consistency may underperform in all but two cases.

► **Different clients may want different consistencies even when accessing**

the same data. Often, systems bind a specific consistency to a particular dataset or class of data. For example, it is generally assumed that bank data must be strongly consistent while shopping cart data needs only eventually consistency. The baseball example shows that the desired consistency depends as much on who is reading the data as on the type of data.

► **Even simple databases may have diverse users with different consistency needs.** A baseball score is one of the simplest databases imaginable, consisting of only two numbers. Nevertheless, it effectively illustrates the value of different consistency options.

► **Clients should be able to choose their desired consistency.** The system cannot possibly predict or determine the consistency that is required by a given application or client. The preferred consistency often depends on how the data is being used. Moreover, knowledge of who writes data or when data was last written can sometimes allow clients to perform a relaxed consistency read, and obtain the associated benefits, while reading up-to-date data.

The main argument often expressed against providing eventual consistency is that it increases the burden on application developers. This may be true, but the extra burden need not be excessive. The first step is to define consistency guarantees developers can understand; observe that the six guarantees presented in Table 1 are each described in a few words. By having the storage system perform write operations in a strict order, application developers can avoid the complication of dealing with update conflicts from concurrent writes. This leaves developers with the job of choosing their desired read consistency. This choice requires a deep understanding of the semantics of their application, but need not alter the basic structure of the program. None of the code snippets that were provided in the previous section required any additional lines to deal specifically with stale data.

Cloud storage systems that offer only strong consistency make it easy for developers to write correct programs but may miss out on the benefits of relaxed consistency. The inherent trade-offs between consistency,

performance, and availability are tangible and may become more pronounced with the proliferation of geo-replicated services. This suggests that cloud storage systems should at least consider offering a larger choice of read consistencies. Some cloud providers already offer two both strongly consistent and eventually consistent read operations, but this article shows their eventual consistency model may not be ideal for applications. Allowing cloud storage clients to read from diverse replicas with a choice of several consistency guarantees could benefit a broad class of applications as well as lead to better resource utilization and cost savings. □

References

1. Abadi, D. Consistency tradeoffs in modern distributed database system design. *IEEE Computer*, (Feb. 2012).
2. Amazon. Amazon DynamoDB; <http://aws.amazon.com/dynamodb/>.
3. Anderson, E., Li, X., Shah, M., Tucek, J. and Wylie, J. What consistency does your key-value store actually provide? In *Proceedings of the Usenix Workshop on Hot Topics in Systems Dependability*, (2010).
4. Bailis, P., Venkataraman, S., Franklin, M., Hellerstein, J. and Stoica, I. Probabilistically bounded staleness for practical partial quorums. In *Proceedings VLDB Endowment*, (Aug. 2012).
5. Brewer, E. CAP twelve years later: How the "rules" have changed. *IEEE Computer*, (Feb. 2012).
6. Calder, B. et. al. Windows azure storage: A highly available cloud storage service with strong consistency. In *Proceedings ACM Symposium on Operating Systems Principles*, (Oct. 2011).
7. Cooper, B., Ramakrishnan, R., Srivastava, U., Silberstein, A., Bohannon, P., Jacobsen, H.A., Puz, N., Weaver, D. and Yerneni, R. PNUTS: Yahoo!'s hosted data serving platform. In *Proceedings International Conference on Very Large Data Bases*, (Aug. 2008).
8. Google. Read Consistency & Deadlines: More Control of Your Datastore. Google App Engine Blob, Mar. 2010; <http://googleappengine.blogspot.com/2010/03/read-consistency-deadlines-more-control.html>.
9. Kraska, T., Hentschel, M., Alonso, G. and Kossmann, D. Consistency rationing in the cloud: Pay only when it matters. In *Proceedings International Conference on Very Large Data Bases*, (Aug. 2009).
10. Saito, Y. and Shapiro, M. Optimistic replication. *ACM Computing Surveys*, (Mar. 2005).
11. Terry, D., Demers, A., Petersen, K., Spreitzer, M., Theimer, M., and Welch, B. Session guarantees for weakly consistent replicated data. In *Proceedings IEEE International Conference on Parallel and Distributed Information Systems*, (1994).
12. Vogels, W. Eventually consistent. *Commun. ACM*, (Jan. 2009).
13. Wada, H., Fekete, A., Zhao, L., Lee, K. and Liu, A. Data consistency properties and the trade-offs in commercial cloud storages: The consumers' perspective. In *Proceedings CIDR*, (Jan. 2011).

Doug Terry (terry@microsoft.com) is a Principal Researcher in the Microsoft Research Silicon Valley Lab, Mountain View, CA.